

SAS Club 2025:

Datenmanagement News SAS SpeedyStore & SAS/ACCESS Interface to DuckDB

Wien, 23.10.2025

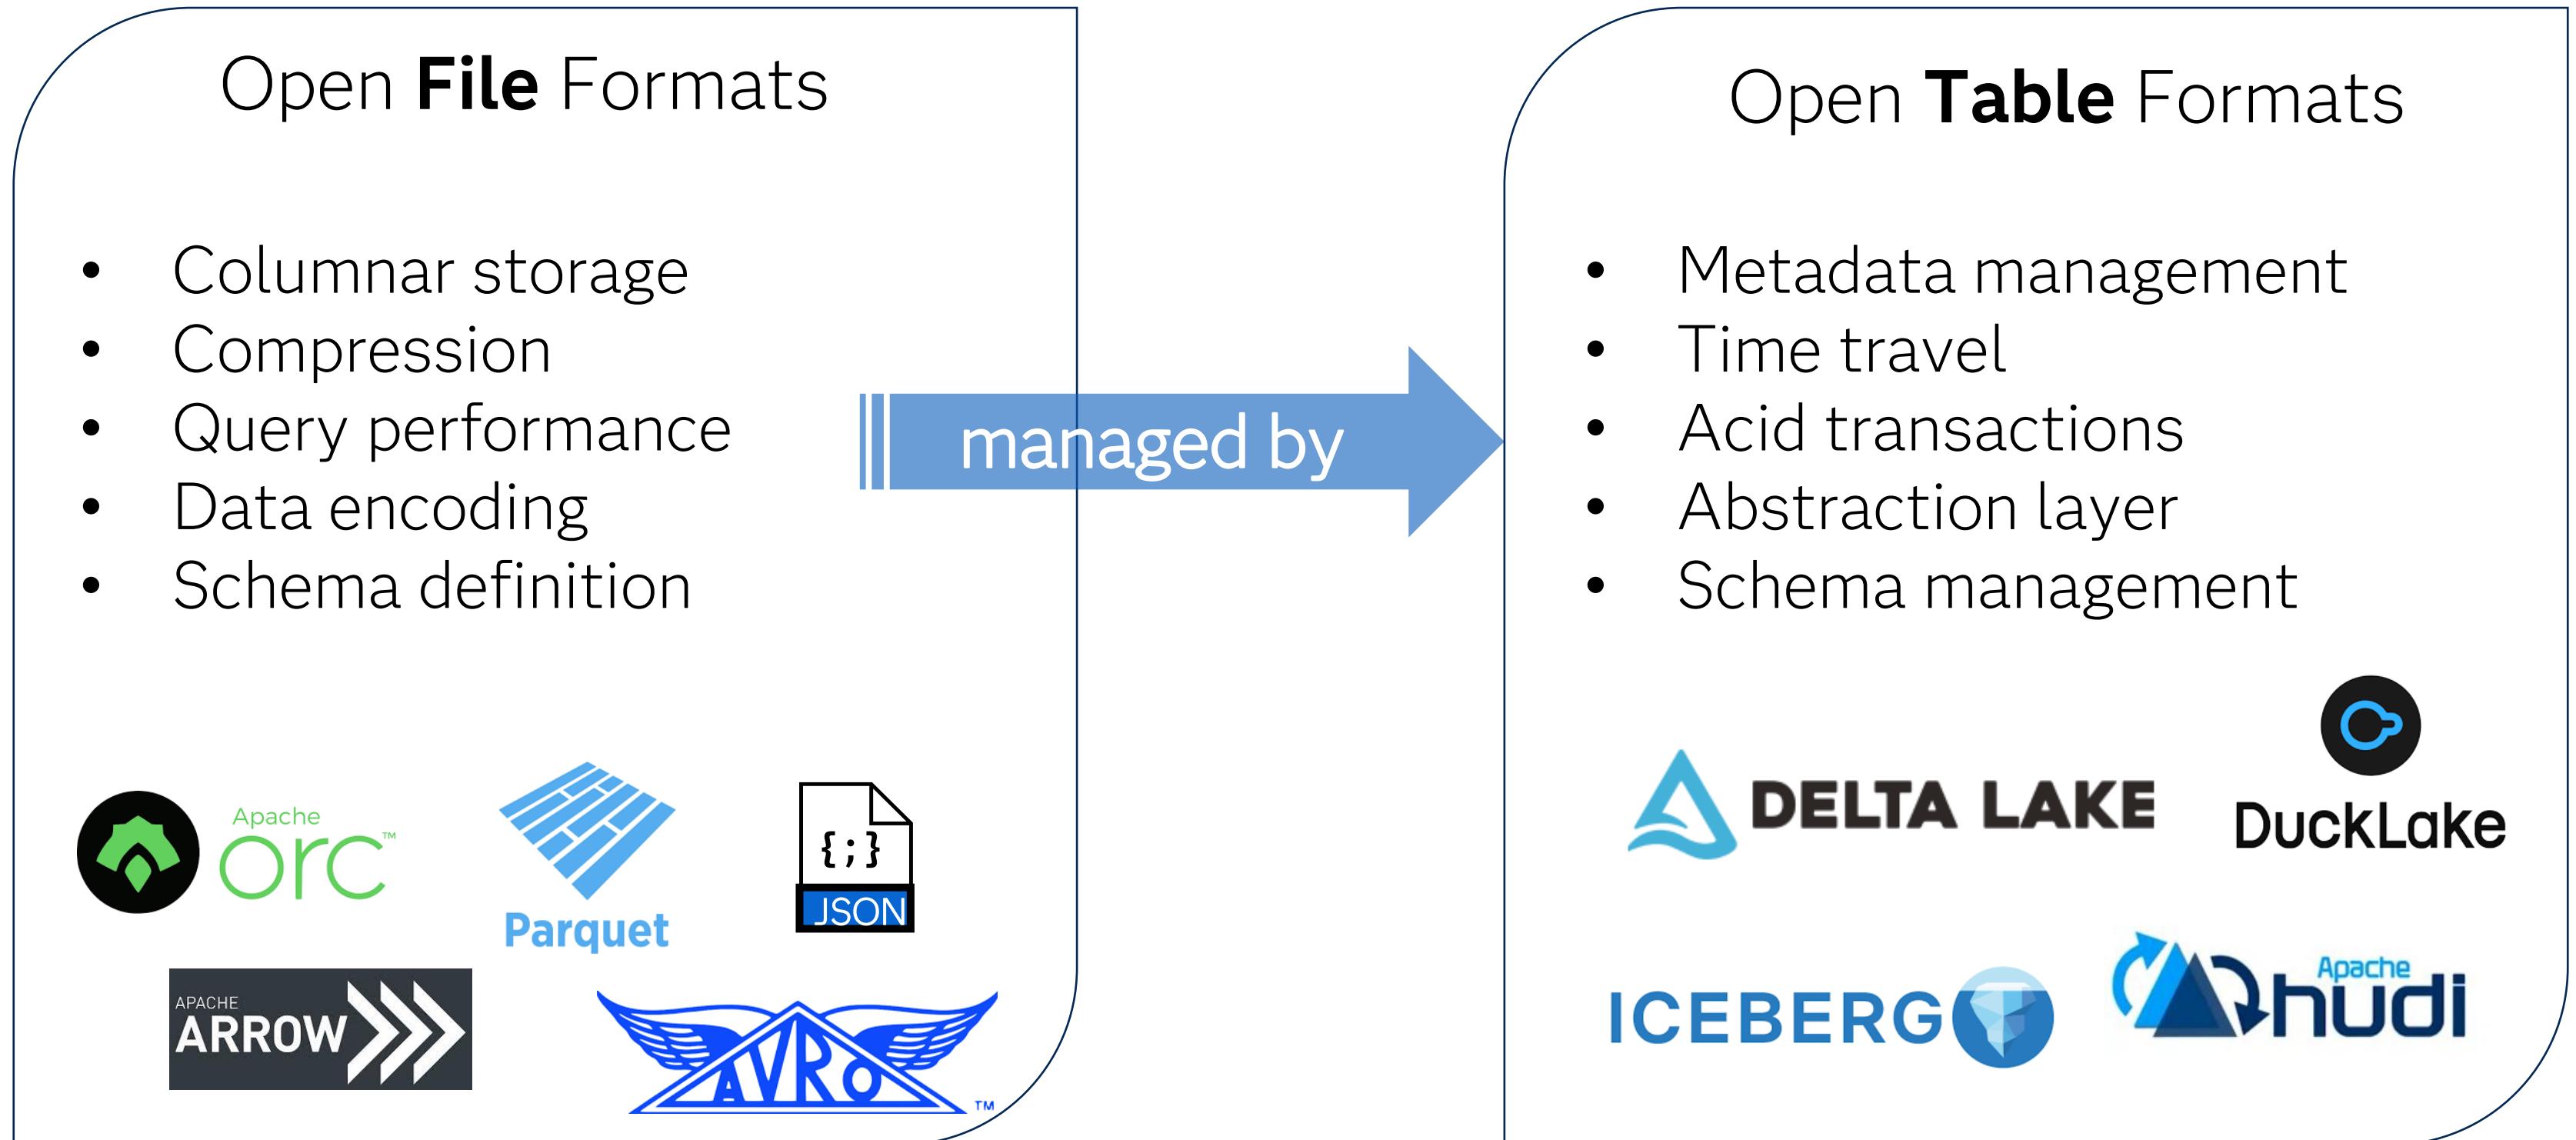


Data Landscape

2025 and beyond

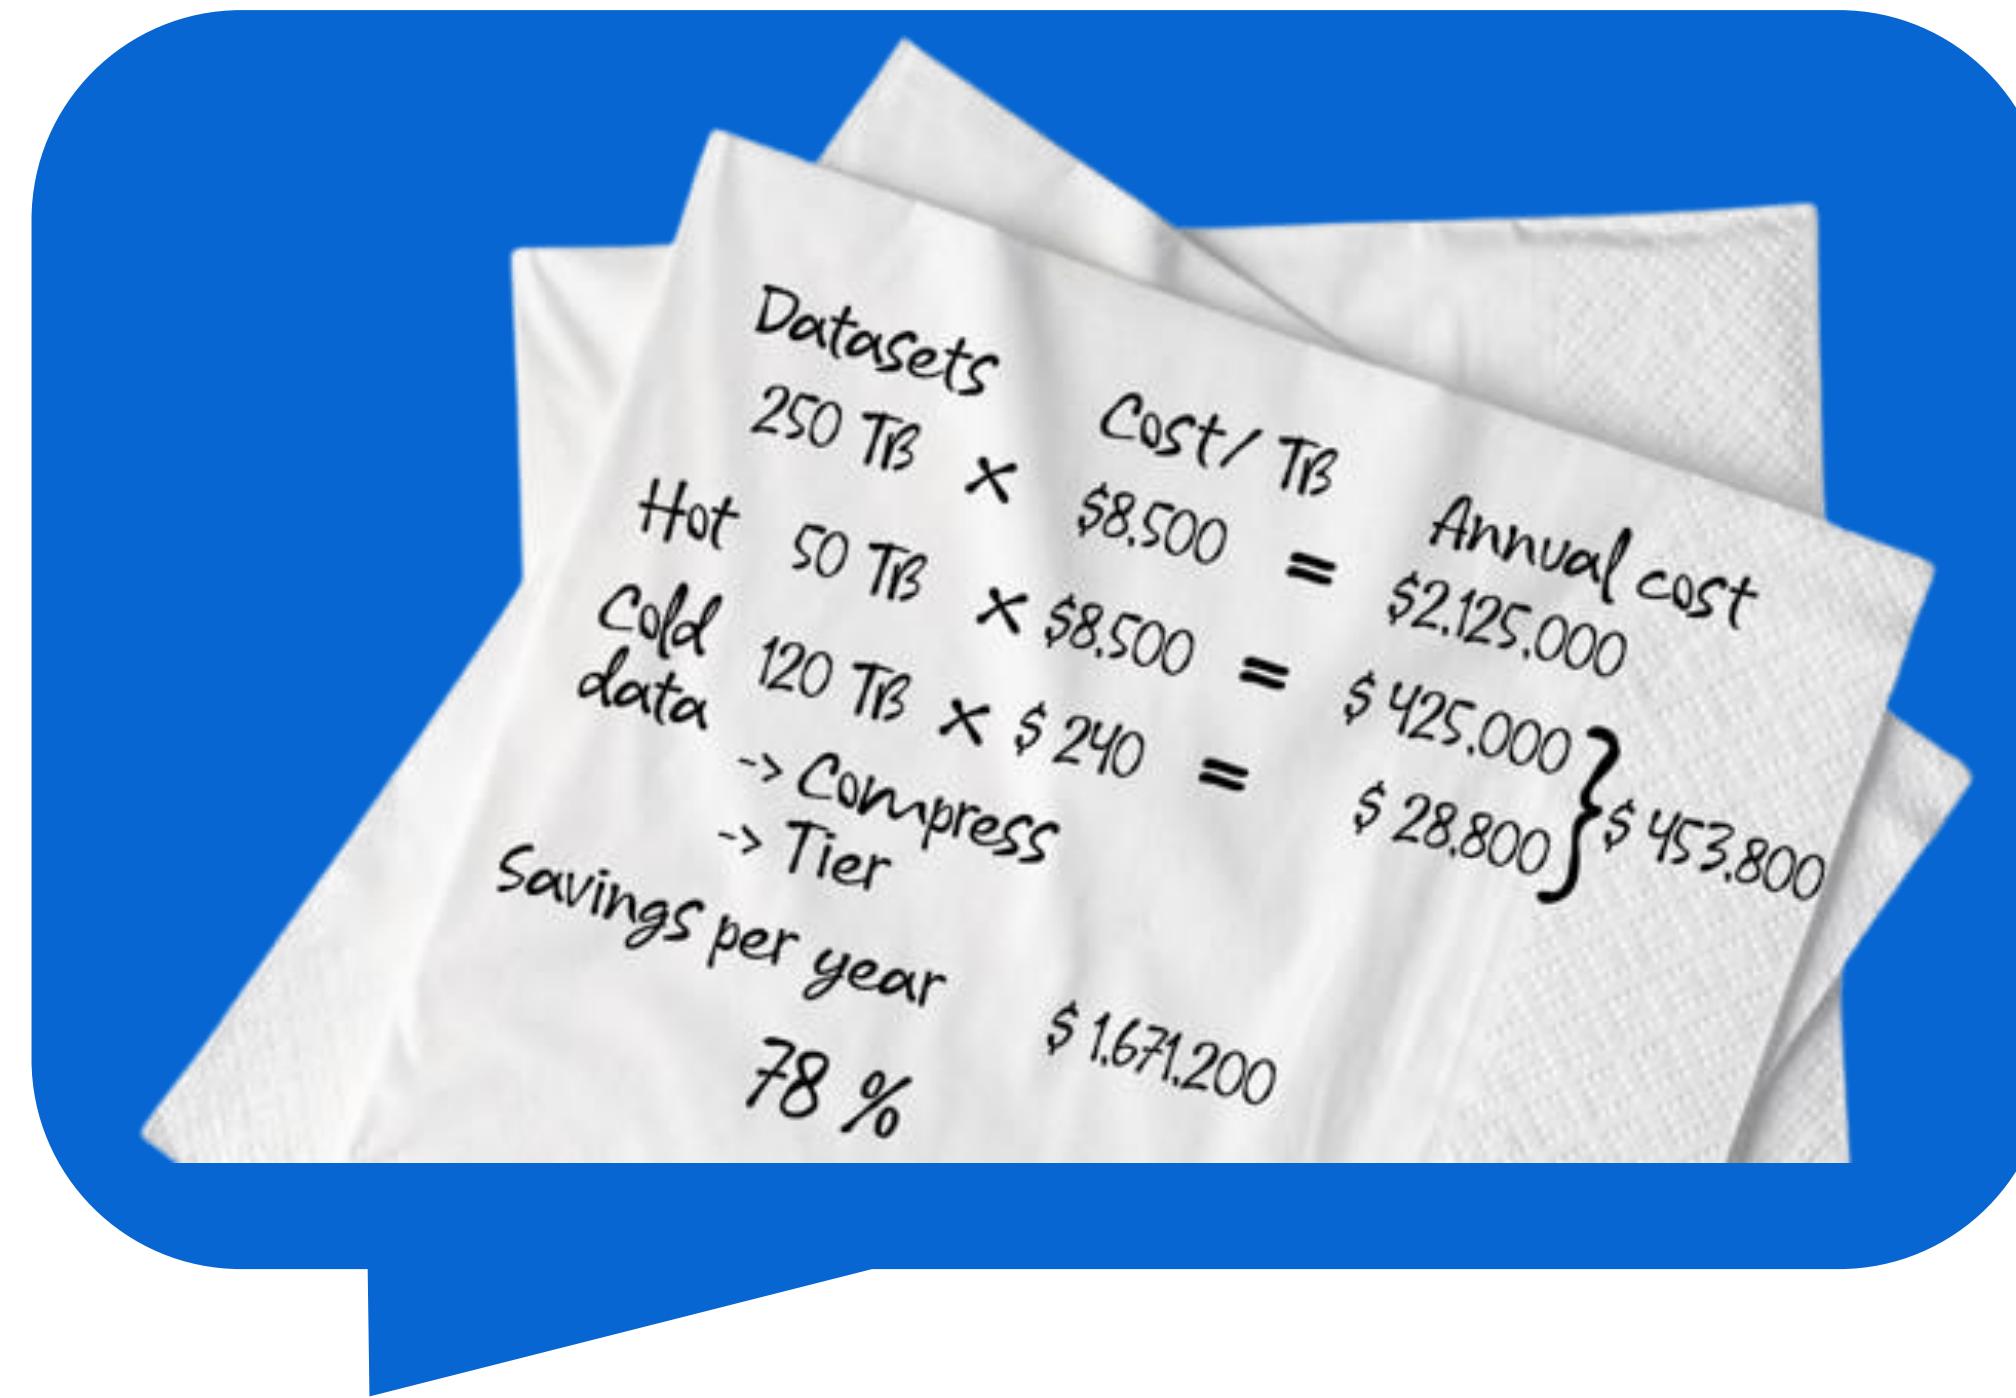


Continued Evolution



Storage Cost Optimization

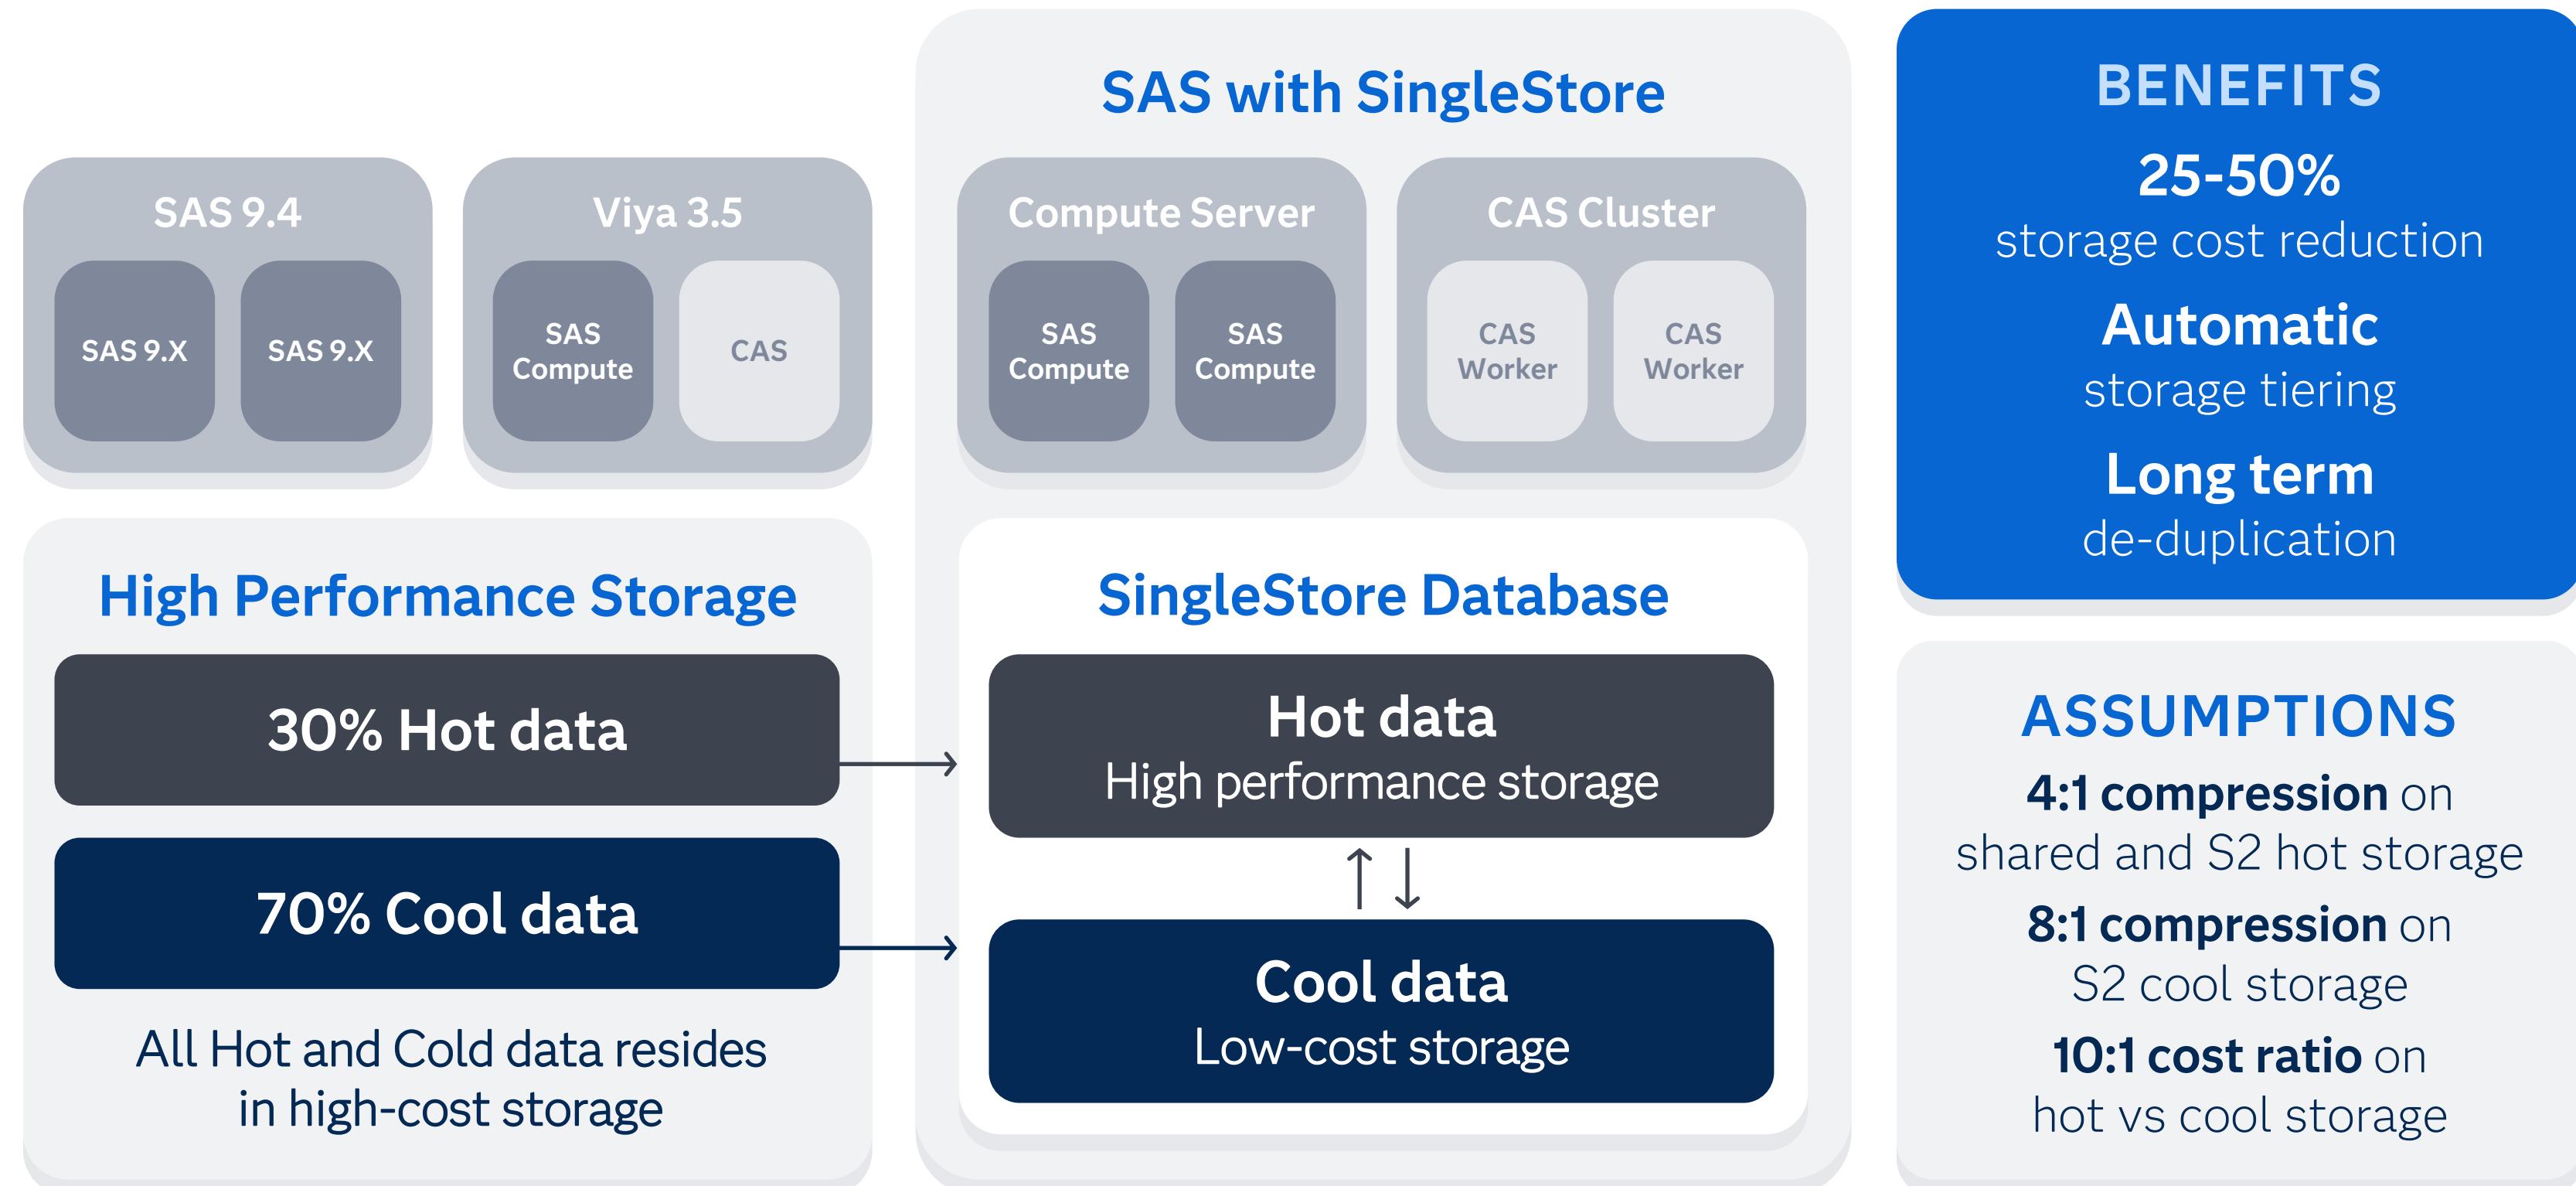
Estimated cost savings with Data Tiering



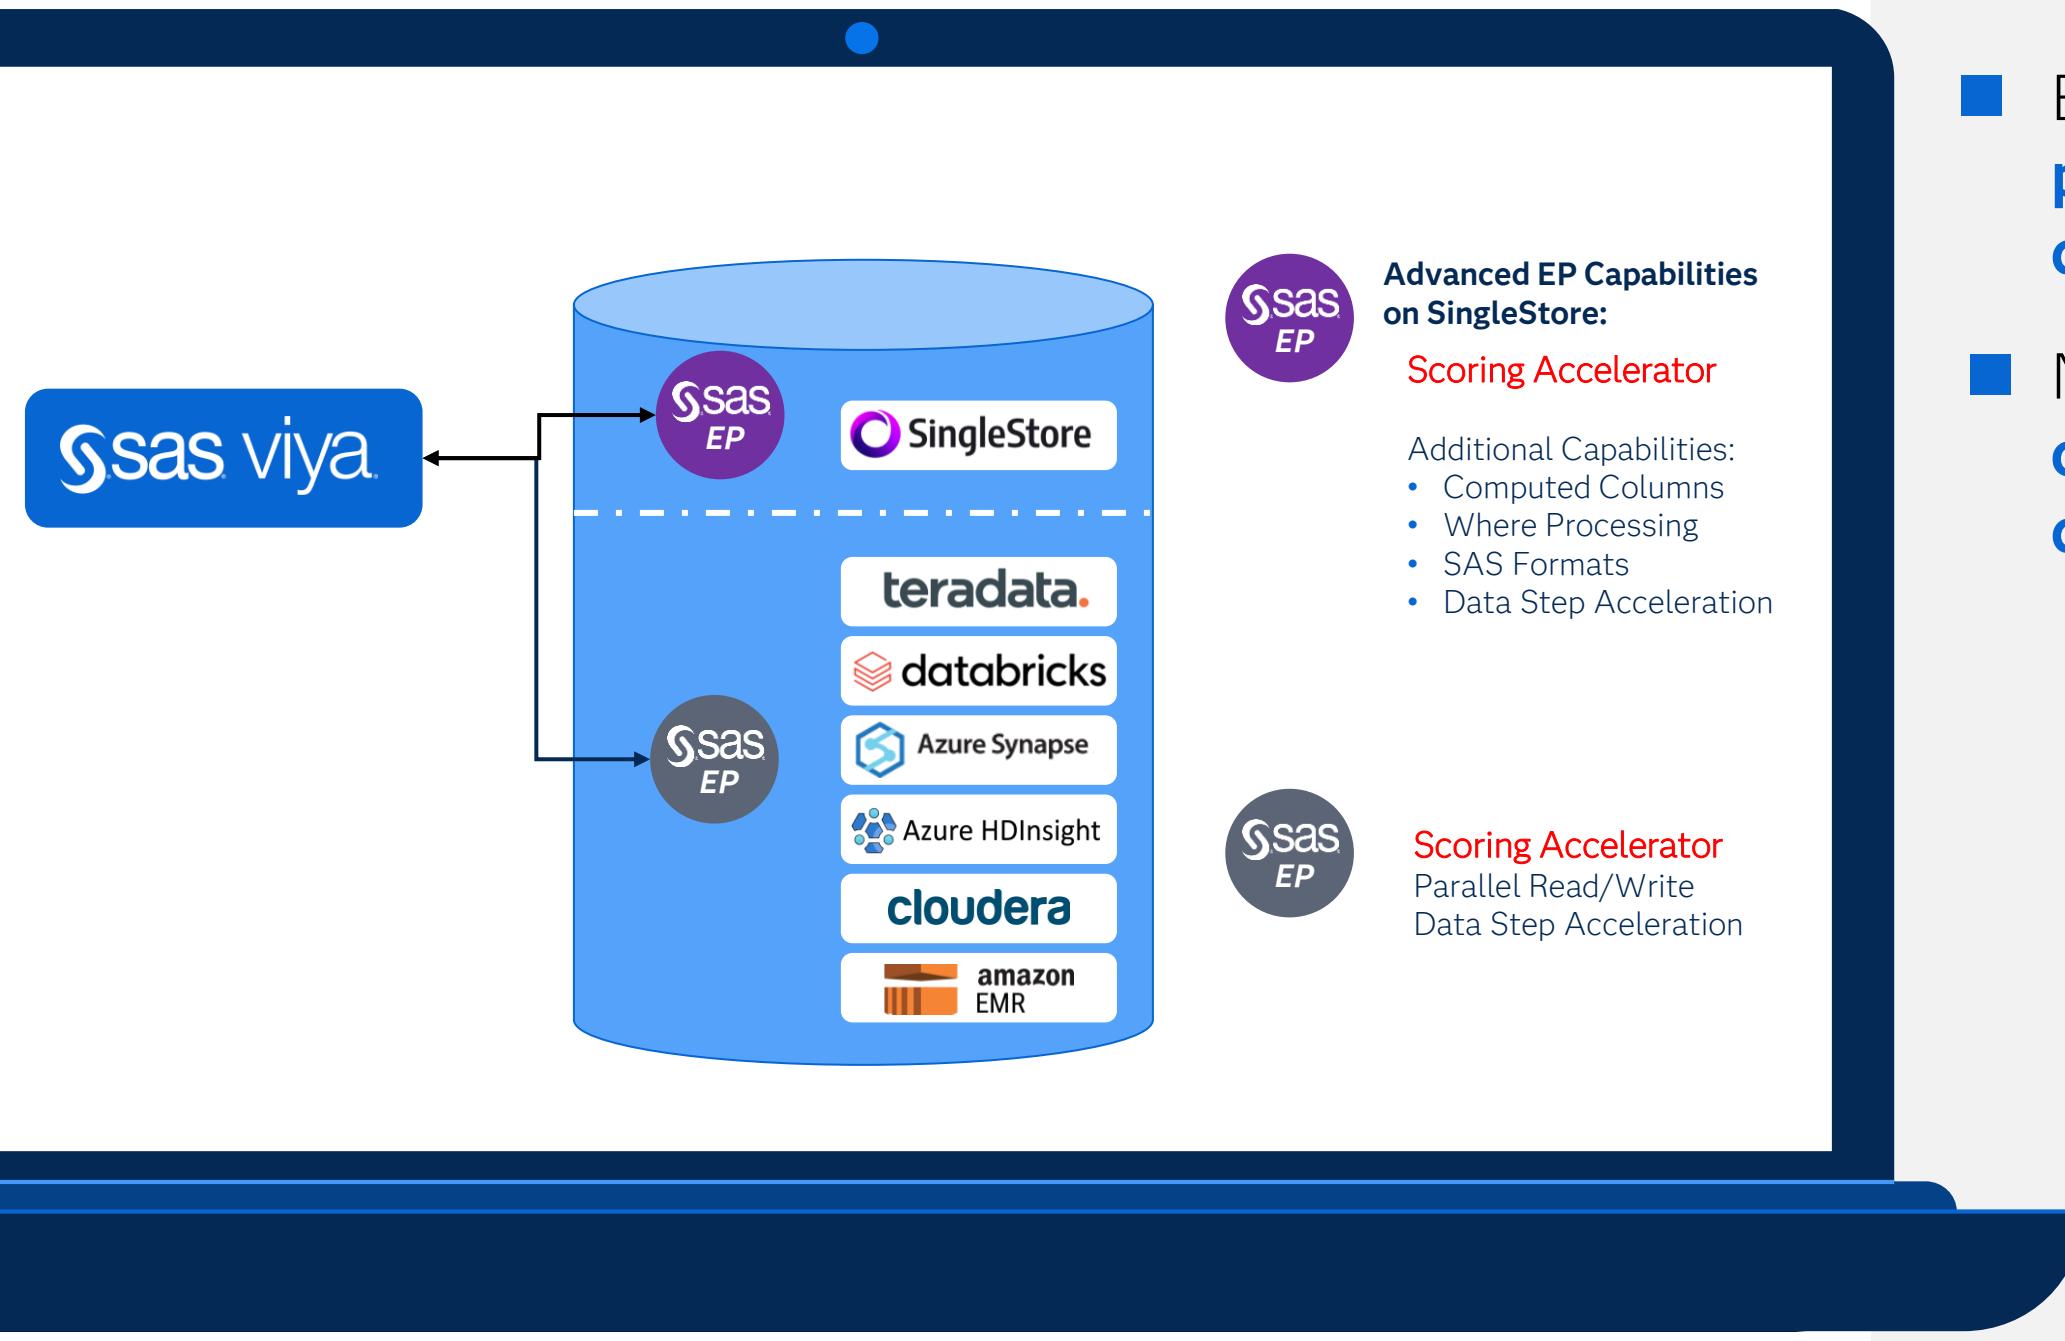
SAS SpeedyStore



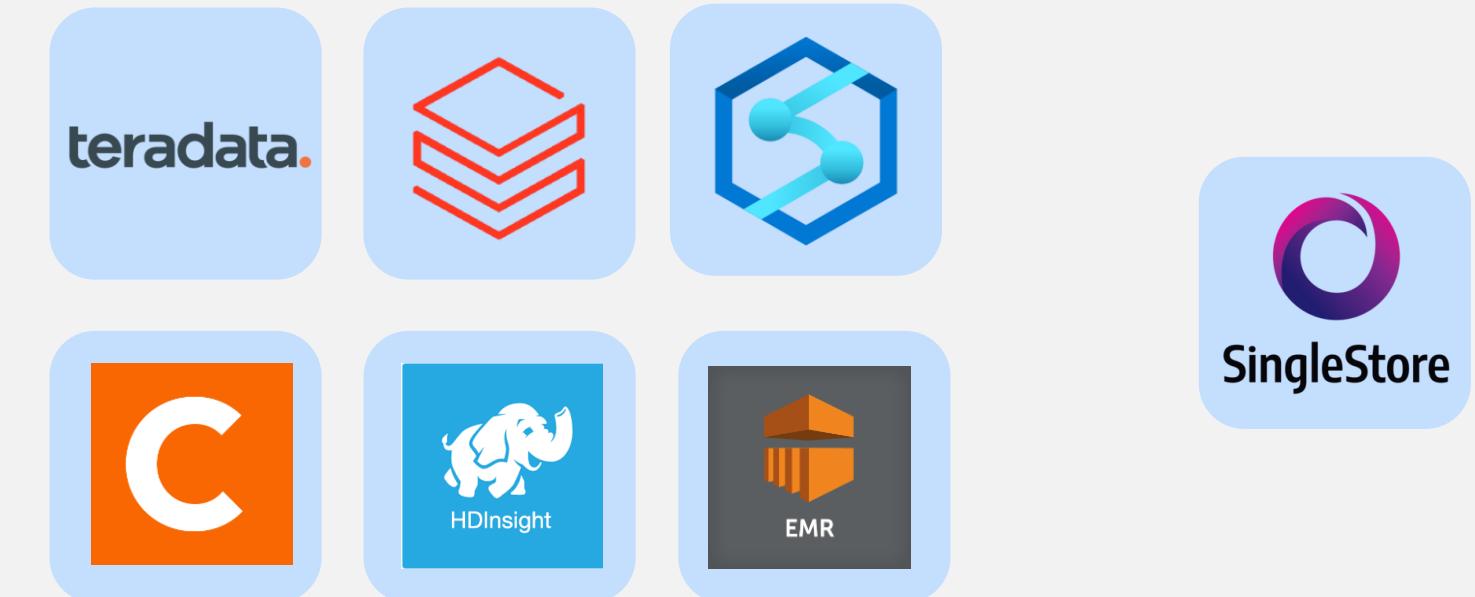
SAS SpeedyStore Data Storage Savings



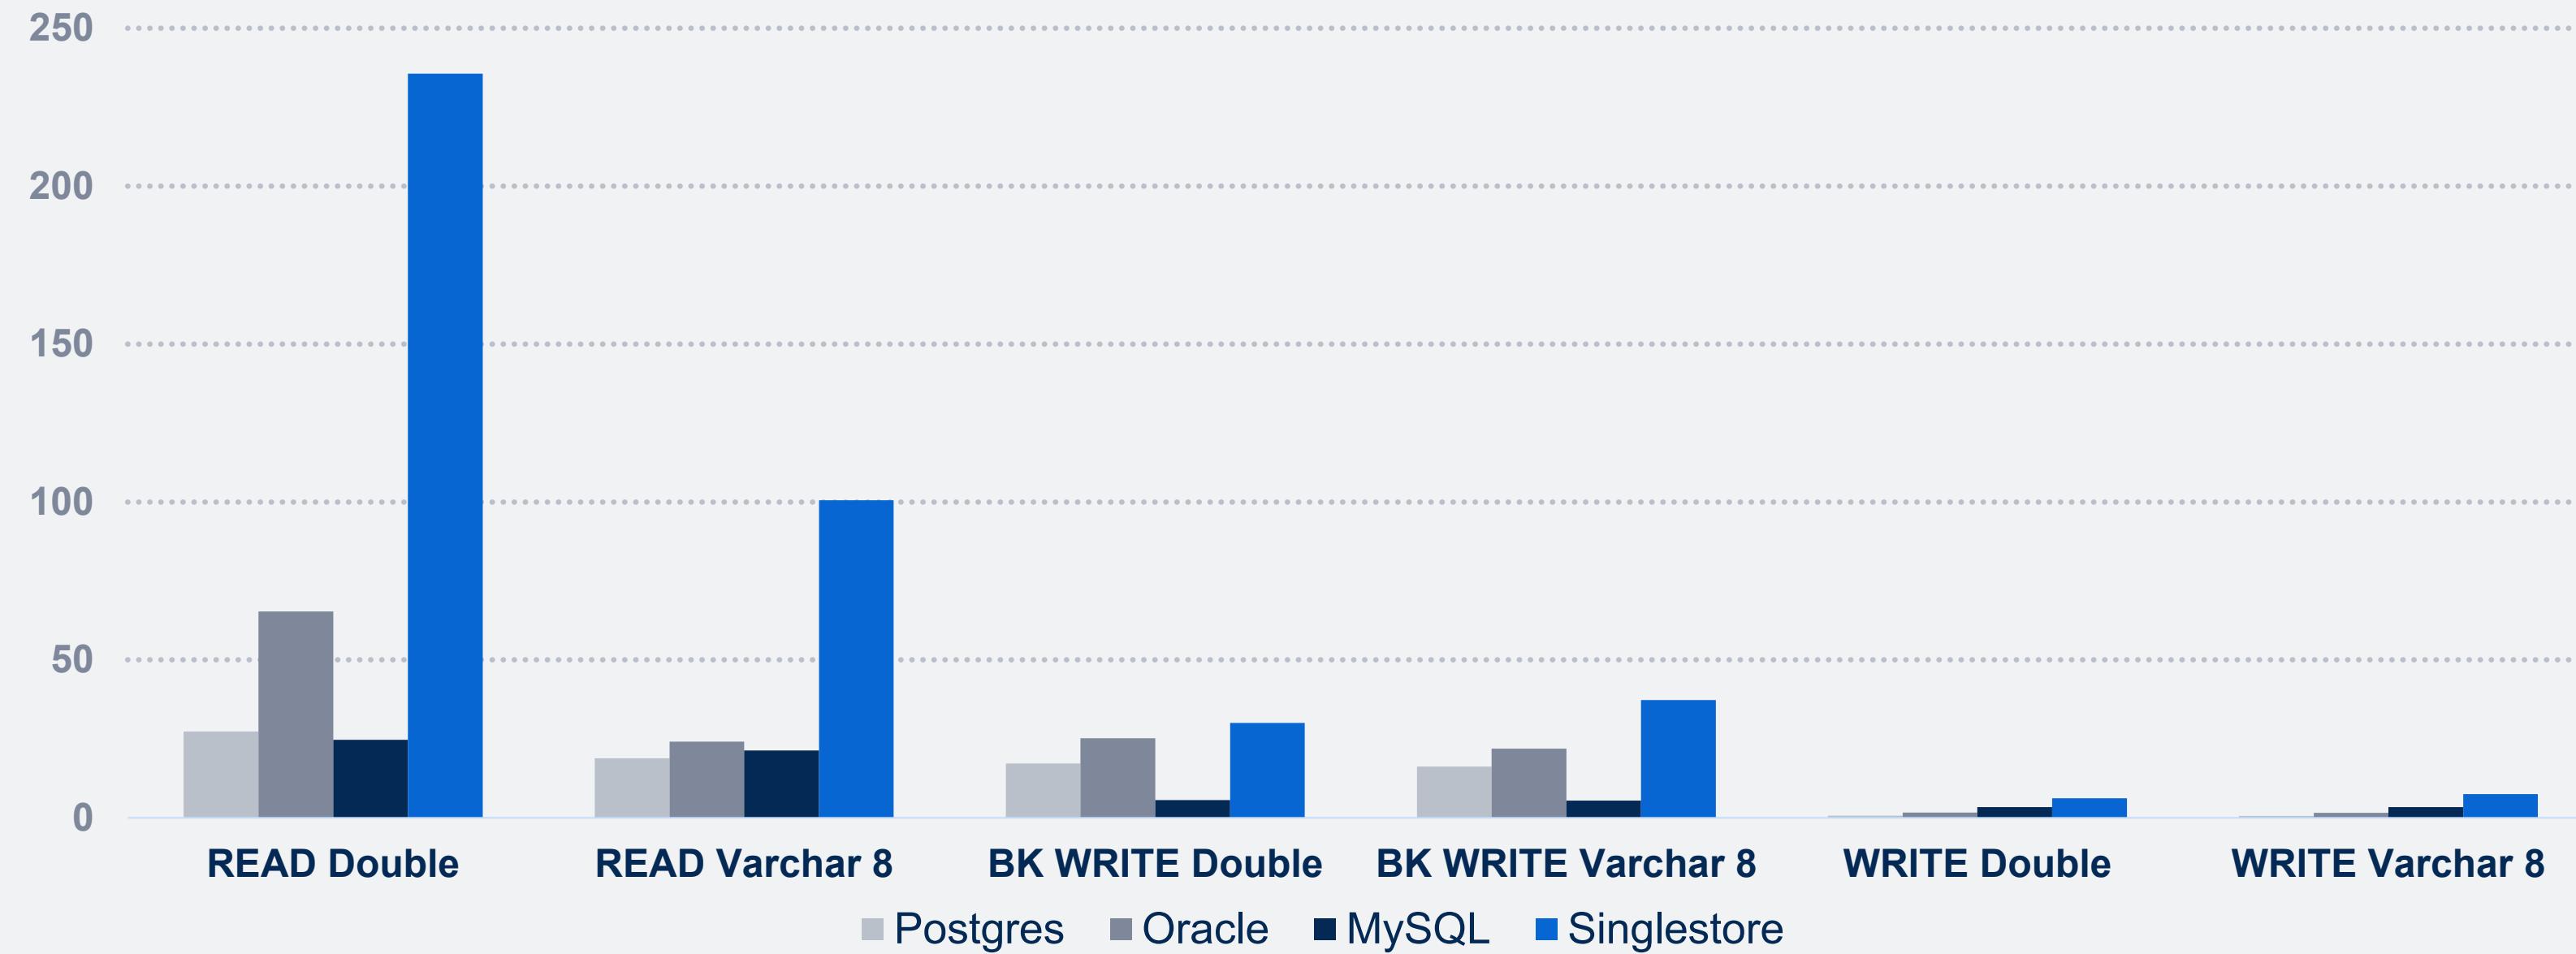
Deploy & Score Models In-Database with Scoring Accelerator



- Efficient way to **process widespread data**
- Moves the **computation to the data**
- Eliminates massive data movement & **reduce processing times**
- Improves data security as **data never leaves the cluster**



Performance in MBps



SAS Viya Performance

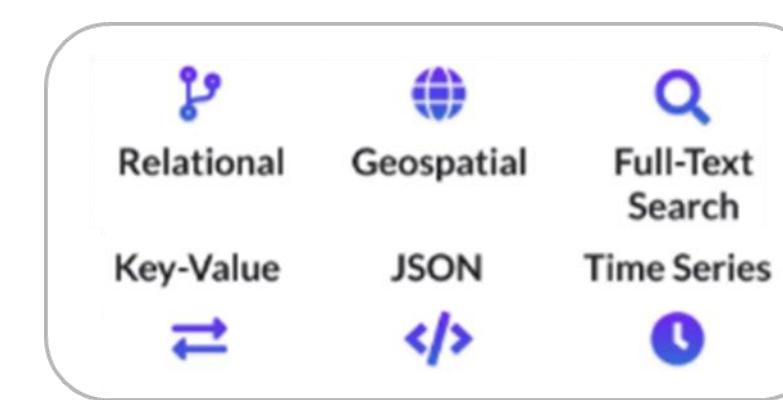
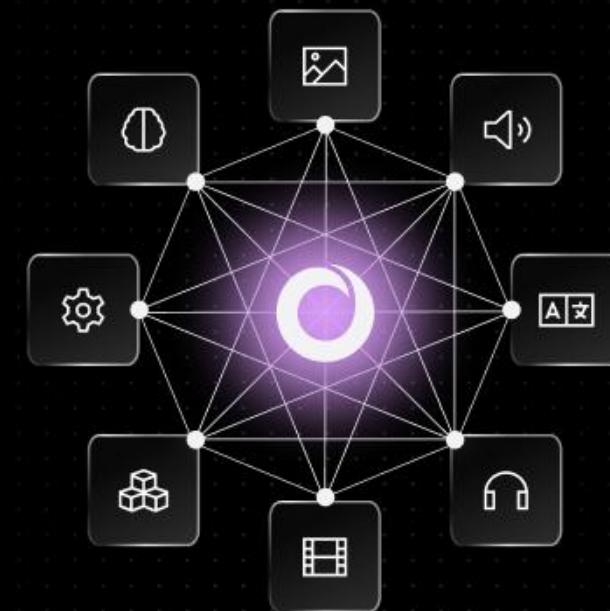
READ 10 million rows

BULKLOAD WRITE 2.5 million rows

WRITE 100k rows

Support JSON, key-value, time series, vectors

SpeedyStore supports all data types



BENEFITS

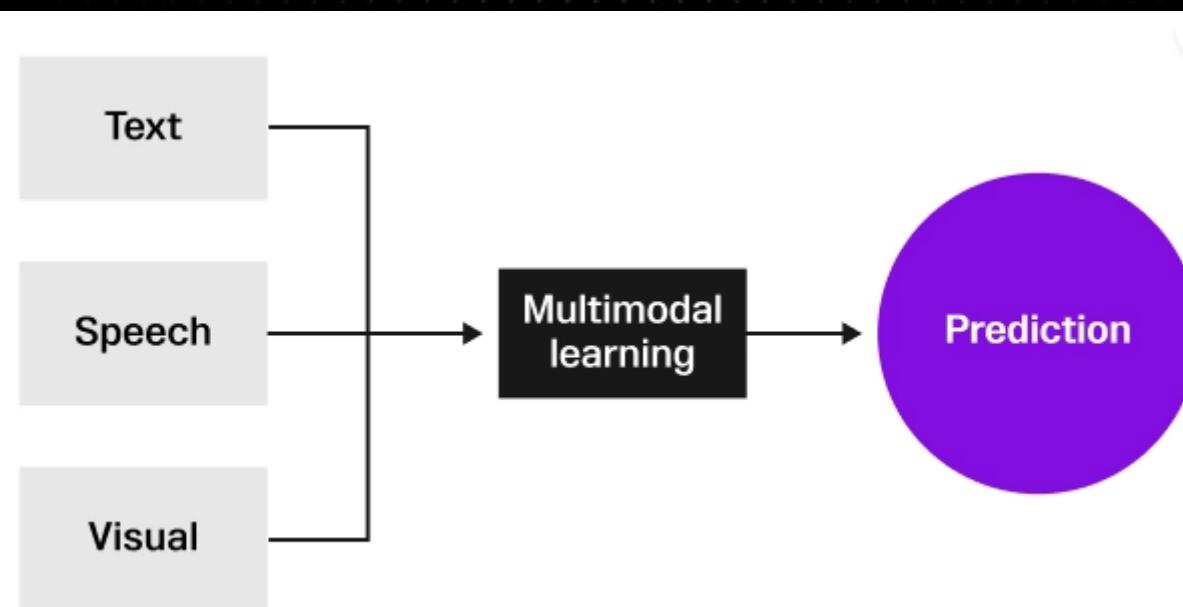
Consolidate your database landscape, only one database to manage

Get access to all data in one database

Reduction and simplification of data pipelines

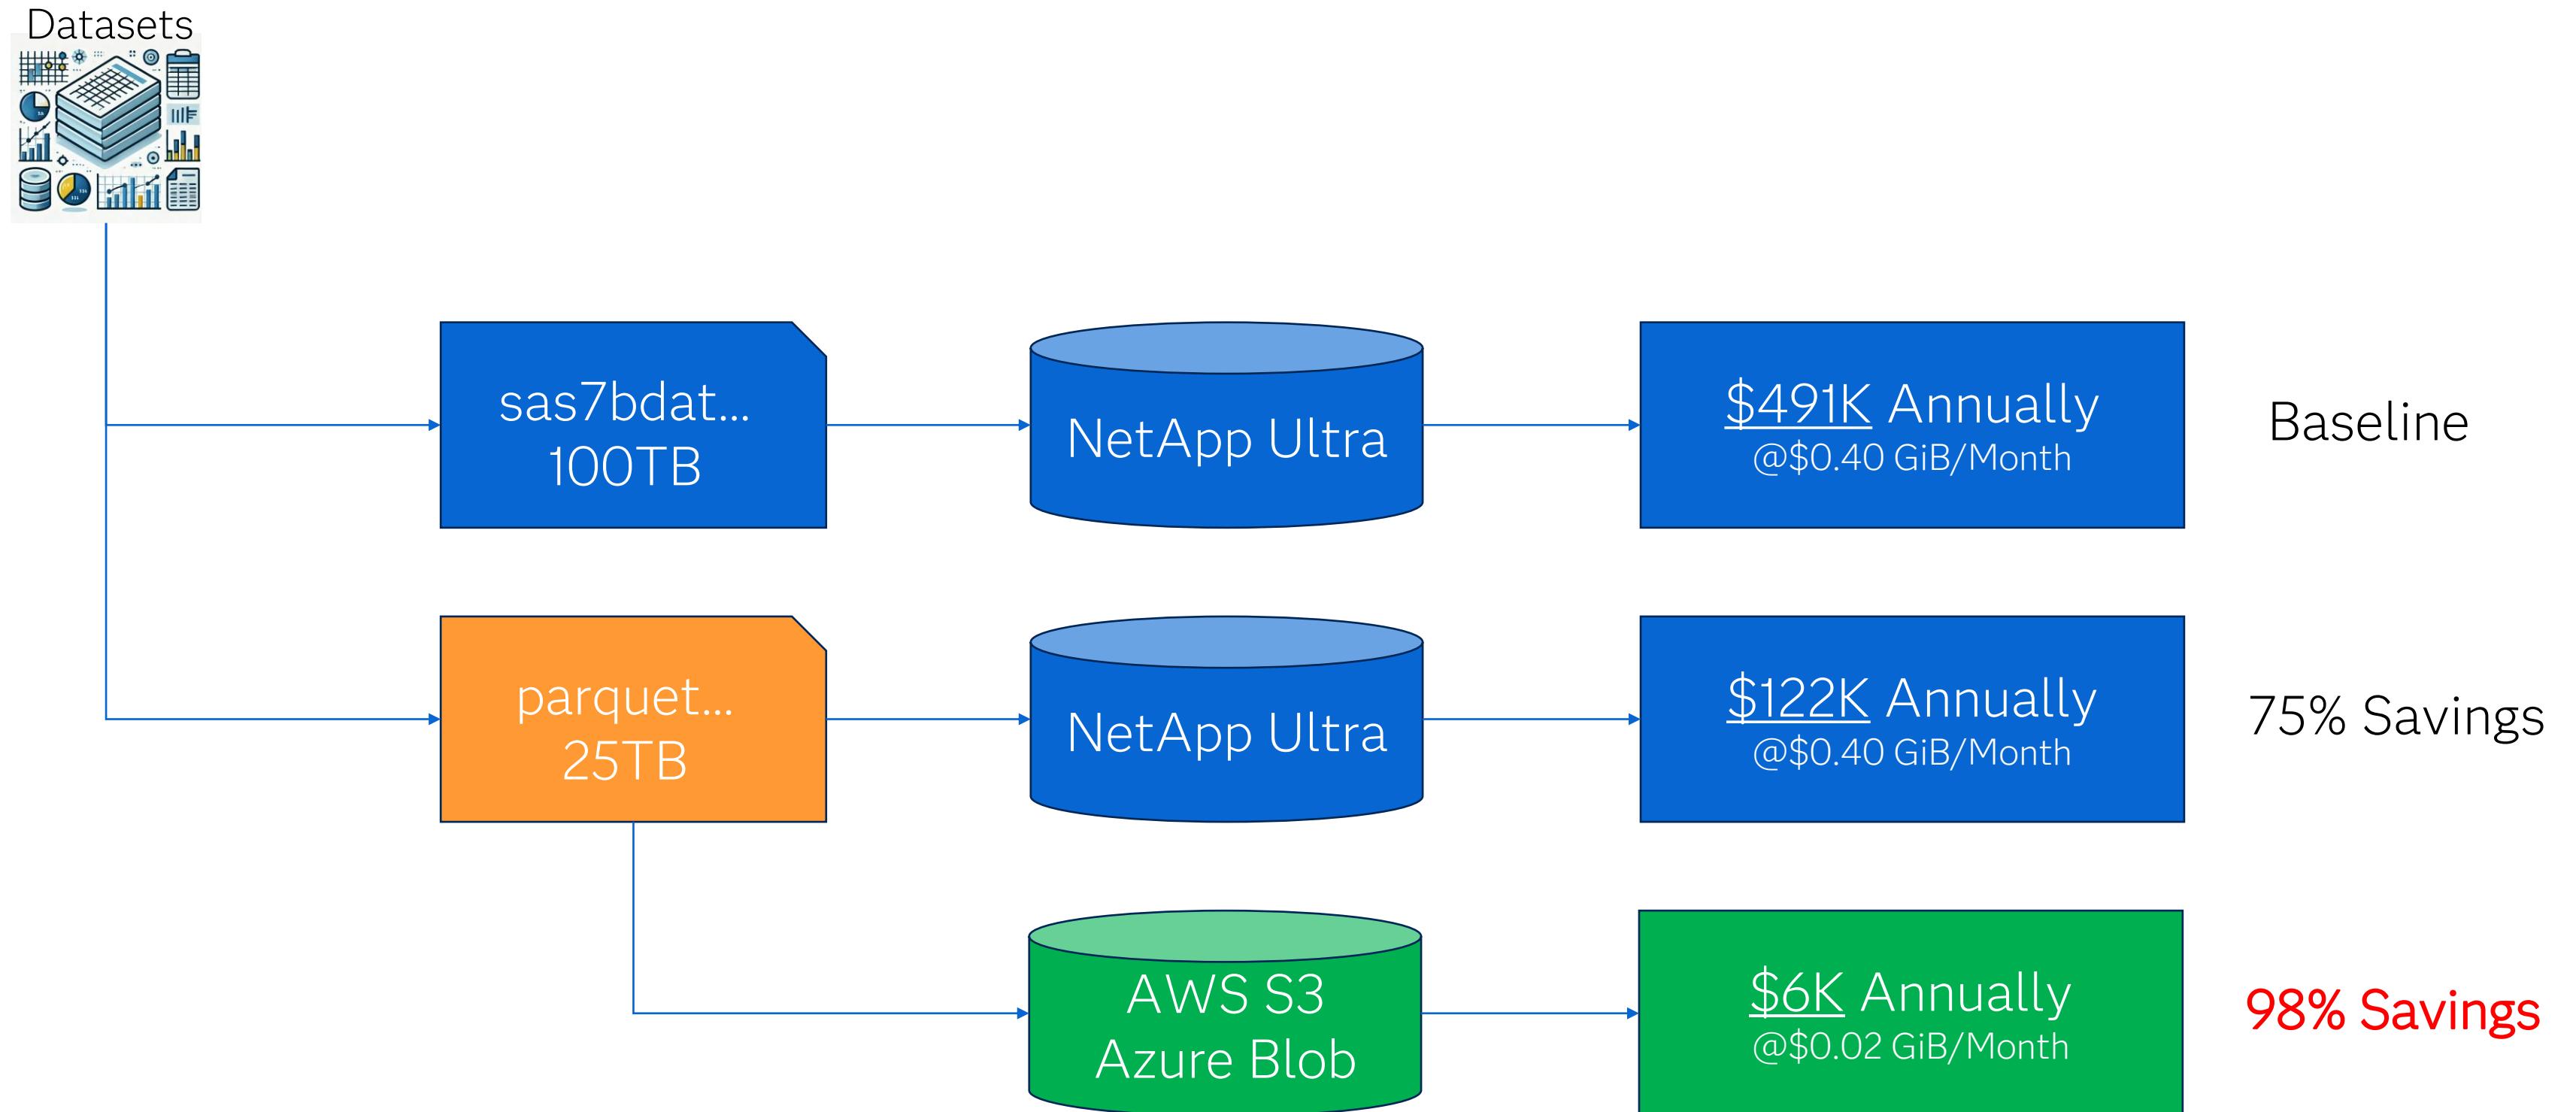
Speeds up the development process

Easy to combine different data types to build multi modal models and applications



Storage Cost Optimization

Example



SAS/Access Interface to DuckDB



 DuckDB is an open source, in-process analytics engine

 Vectorised, optimised query execution on columnar storage

 Requires no server installation and no external dependencies

 Included as part of SAS Viya from 2025.07 release

 Extensive support for open file formats and **open table formats**

SAS/Access Interface to DuckDB

Performance (27GB dataset)

```
PROC SQL;
  SELECT
    passenger_count,
    payment_type,
    count(*)      AS num_trips,
    avg(trip_distance) AS avg_distance,
    avg(fare_amount) AS avg_fare,
    avg(tip_amount) AS avg_tip
  FROM
    saslib.yellow_tripdata_2011
  WHERE
    passenger_count is not NULL AND
    passenger_count > 0 AND
    passenger_count < 5 AND
    trip_distance < 100 AND
    trip_distance > 0
  GROUP BY
    passenger_count, payment_type
  ORDER BY
    payment_type, passenger_count;
QUIT;
```

NOTE: PROCEDURE SQL used (Total process time):
real time 1:56.27
cpu time 1:22.98

```
PROC SQL;
  SELECT
    passenger_count,
    payment_type,
    count(*)      AS num_trips,
    avg(trip_distance) AS avg_distance,
    avg(fare_amount) AS avg_fare,
    avg(tip_amount) AS avg_tip
  FROM
    duklib.'2011/*.parquet'n
  WHERE
    passenger_count is not NULL AND
    passenger_count > 0 AND
    passenger_count < 5 AND
    trip_distance < 100 AND
    trip_distance > 0
  GROUP BY
    passenger_count, payment_type
  ORDER BY
    payment_type, passenger_count;
QUIT;
```

NOTE: PROCEDURE SQL used (Total process time):
real time 2.54 seconds
cpu time 26.24 seconds

SASIODUK vs BASE engine

